
Creating Docker Images



olsen software

Getting Demos and Lab Code

- We've provided various demos and labs here:
 - <http://olsensoft.com/DockerOpenShift/DoshDev.zip>
- You must download the zip into your Linux box, and unzip
 - To do this, run the following commands in your Linux box:

```
wget http://olsensoft.com/DockerOpenShift/DoshDev.zip -O temp.zip; \  
unzip temp.zip; \  
rm temp.zip
```

Contents

1. Creating a simple Docker image
2. A closer look at Dockerfile instructions
3. Parameterizing Docker containers

Annex

- The layered filesystem of Docker images

1. Creating a Simple Docker Image

- Overview
- Defining a Dockerfile
- Understanding the Dockerfile
- Our Python program
- Building the image
- Viewing images in the local Docker registry
- Running a container

Overview

- In this section, we're going to show how to create a simple Docker image
 - Remember, a Docker image is a "black box" executable package
 - It includes everything needed to run an application
- In our case, we'll create a Docker image containing:
 - The Python runtime
 - Our Python program
 - A command to execute the Python program
- For the demo code, see this folder:
 - `PythonImage1`

Defining a Dockerfile

- In order to define a Docker image, you define a special file named `Dockerfile` (by default)
 - Specifies build instructions, telling Docker how to build an image
- Here's the Dockerfile in the demo folder, `PythonImage1`

```
FROM python:3.8.5  
  
COPY ./main.py app.py  
  
ENTRYPOINT ["python", "app.py"]
```

- See following slides for an explanation
 - Also see <https://docs.docker.com/engine/reference/builder/>

Understanding the Dockerfile (1 of 3)

```
FROM python:3.8.5
```

```
COPY ./main.py app.py
```

```
ENTRYPOINT ["python", "app.py"]
```

- A Dockerfile must start with a FROM instruction
 - Specifies the "base image" from which you are building
 - Our image is based on Python version 3.8.5
 - If you want to create an image from scratch, use FROM SCRATCH
- When you run this Dockerfile to build your image...
 - Docker will see if you've already downloaded Python into your local Docker registry
 - If not, Docker will automatically download it from Docker Hub

Understanding the Dockerfile (2 of 3)

```
FROM python:3.8.5  
COPY ./main.py app.py  
ENTRYPOINT ["python", "app.py"]
```

- A Dockerfile specifies files and folders to copy into the Docker image
 - Use the COPY instruction to copy files and folders into the image
 - In our example, we copy `main.py` into the image
 - Inside the image, the Python file will be named `app.py`

Understanding the Dockerfile (3 of 3)

```
FROM python:3.8.5
COPY ./main.py app.py
ENTRYPOINT ["python", "app.py"]
```

- The ENTRYPOINT instruction specifies what to actually execute inside the Docker image
 - In our example, we run the Python interpreter and tell it to run our Python program
 - We've used the "exec form" of the ENTRYPOINT instruction here, which is JSON array syntax
- The ENTRYPOINT instruction also supports "shell form" syntax, and executes it in `/bin/sh -c`

```
ENTRYPOINT python app.py
```

Our Python Program

- Here's our Python program 😊

```
print ("Hello from my uber cool Python app #1")
```

Building the Image (1 of 2)

- Use the `docker build` command to build a Docker image

```
docker build -t mypython1 .
```

- `-t`
 - Specifies the tag name for the image
- `.`
 - Specifies the current directory as the build context
 - You can also specify a path, URL, or tarball
- Note:
 - Assumes the dockerfile is called `Dockerfile` (use `-f` if different)

Building the Image (2 of 2)

- Here's how it looks when we use `docker build`

```
docker build -t mypython1 .
```



```
$ docker build -t mypython1 .
Sending build context to Docker daemon 3.072kB
Step 1/3 : FROM python:3.8.5
3.8.5: Pulling from library/python
d6ff36c9ec48: Pull complete
c958d65b3090: Pull complete
edaf0a6b092f: Pull complete
80931cf68816: Pull complete
7dc5581457b1: Pull complete
87013dc371d5: Pull complete
ddb5b2d86fe3: Pull complete
4cb6f1e38c2d: Pull complete
c2df8846f270: Pull complete
Digest: sha256:bc765f71aaa90648de6cfa356ec201d50549031a244f48f8f477f386517c5d1b
Status: Downloaded newer image for python:3.8.5
--> a7cda474cef4
Step 2/3 : COPY ./main.py app.py
--> b58c54a9abac
Step 3/3 : ENTRYPOINT ["python", "app.py"]
--> Running in 0d17e4db3be2
Removing intermediate container 0d17e4db3be2
--> 465145c39144
Successfully built 465145c39144
Successfully tagged mypython1:latest
```

Viewing Images in the Local Docker Registry

- You can view images in the Docker registry

```
docker image ls
```



```
$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
mypython1	latest	465145c39144	3 minutes ago	882MB
python	3.8.5	a7cda474cef4	4 days ago	882MB

Running a Container

- You can run containers based on the image

```
docker run mypython1
```



```
$ docker run mypython1  
Hello from my uber cool Python app #1
```

2. A Closer Look at Dockerfile Instructions

- Overview
- The ARG instruction
- The RUN instruction
- The WORKDIR instruction
- The COPY and ADD instructions
- The EXPOSE instruction
- The ENTRYPOINT instruction

Overview

- In this section we're going to look at a more interesting Dockerfile, to explore additional Docker techniques

```
FROM python:3.8.5

ARG PYTHON_MAIN_FILE

RUN mkdir /app

WORKDIR /app

COPY ./requirements.txt /app
COPY ${PYTHON_MAIN_FILE} /app/main.py

RUN pip install --trusted-host pypi.python.org -r requirements.txt

EXPOSE 5000

ENTRYPOINT ["python", "main.py"]
```

- For the demo code, see this folder:
 - PythonImage2

The ARG Instruction

- The ARG instruction allows you to specify build arguments for the Dockerfile, to make it more flexible

```
...  
ARG PYTHON_MAIN_FILE  
...  
COPY ${PYTHON_MAIN_FILE} /app/main.py  
...
```

- You pass argument values into the Dockerfile as part of the Docker build command
 - Via the `--build-arg` flag

```
docker build -t mypython2 \  
--build-arg PYTHON_MAIN_FILE=server.py \  
.
```

The RUN Instruction

- The RUN instruction allows you to run any valid Linux command, e.g. to make a directory in the image

```
RUN mkdir /app
```

- Remember each statement in a Dockerfile creates a separate layer in the union filesystem
 - So it's best to combine multiple Linux commands into a single RUN statement, to minimize the number of layers

```
RUN mkdir /app && cd /app
```



```
RUN mkdir /app  
RUN cd /app
```



The WORKDIR Instruction

- The WORKDIR instruction sets the working directory inside the image
 - All activity that happens inside the image after this statement will use this as the working directory

```
WORKDIR /app
```

- Note that WORKDIR is different to cd
 - WORKDIR sets the context across layers in the image filesystem
 - But cd only applies to the current layer, not subsequent layers

```
WORKDIR /app  
RUN touch logfile.txt      # Creates /app/logfile.txt
```



```
RUN cd /app  
RUN touch logfile.txt      # Creates /logfile.txt
```



The COPY and ADD Instructions (1 of 2)

- The COPY instruction copies files and folders from the host filesystem into the image filesystem

```
COPY ./requirements.txt /app  
COPY ${PYTHON_MAIN_FILE} /app/main.py
```

- The ADD instruction is similar, but more flexible
 - Allows you to specify URLs and tarballs
 - As well as files and folders, of course

```
ADD mytarball.tar /app  
ADD http://olsensoft.com/mydata /data
```

The COPY and ADD Instructions (2 of 2)

- It's common practice to copy all of the contents of a folder on the host system into the image

```
COPY . /app
```

- You can define a `.dockerignore` file to tell Docker to ignore some files and folders during the copying

```
.git  
  
**/*.class  
  
*.md  
!README*.md
```

`.dockerignore`

The EXPOSE Instruction

- The EXPOSE instruction tells Docker that the container listens on the specified network port(s) at runtime

```
EXPOSE 5000
```

- Note: EXPOSE doesn't actually publish the port
 - It just acts as documentation between the person building the image and the person running the container
 - To publish the port when running the container, use the `-p` flag

```
docker run -d -p 5000:5000 mypython2
```

The ENTRYPOINT Instruction

- The ENTRYPOINT instruction tells Docker what to execute in the running container
- Here's the "exec form" of ENTRYPOINT
 - This is generally preferred

```
ENTRYPOINT ["python", "main.py"]
```

- There's also the "shell form" of ENTRYPOINT
 - This executes in a shell, via `/bin/sh -c`

```
ENTRYPOINT python main.py
```

Pinging the Python Rest Service

- The containerized Python application is running a REST service, with endpoints such as `api/Products`
- You can ping the REST endpoint as follows:

```
curl localhost:5000/api/Products
```



```
$ curl localhost:5000/api/Products
[{"id": 0, "description": "Swansea City shirt", "price": 55, "unitsInStock": 500}, {"id": 1, "description": "Cardiff City shirt", "price": 1.99, "unitsInStock": 20000}, {"id": 2, "description": "Bugatti Divo", "price": 4000000, "unitsInStock": 2}, {"id": 3, "description": "Carving Skis", "price": 350, "unitsInStock": 75}, {"id": 4, "description": "Ski Boots", "price": 150, "unitsInStock": 150}, {"id": 5, "description": "55in OLED HDTV", "price": 1800, "unitsInStock": 100}]
```


3. Parameterizing Docker Containers

- The CMD instruction
- Running a container with default arguments
- Running a container and overriding arguments

The CMD Instruction

- The CMD instruction is often used alongside ENTRYPOINT
 - ENTRYPOINT specifies the entrypoint to run, plus stable args
 - CMD specifies additional args, which you might want to override when you run a container instance
- Example - see the EntrypointCmdExample folder

```
FROM alpine:latest
ENTRYPOINT ["ping", "127.0.0.1"]
CMD ["-c", "3"]
```

Dockerfile

- Build the image as normal

```
docker build -t entrypoint_cmd .
```

Running a Container with Default Arguments

- If you run the container as follows, it uses the default command-line args as specified by CMD

```
docker run -t entrypoint_cmd
```



```
$ docker run -t entrypoint_cmd
PING 127.0.0.1 (127.0.0.1): 56 data bytes
64 bytes from 127.0.0.1: seq=0 ttl=64 time=0.751 ms
64 bytes from 127.0.0.1: seq=1 ttl=64 time=0.077 ms
64 bytes from 127.0.0.1: seq=2 ttl=64 time=0.106 ms

--- 127.0.0.1 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.077/0.311/0.751 ms
```

Running a Container and Overriding Arguments

- If you run the container as follows, it overrides the command-line args specified by CMD

```
docker run -t entrypoint_cmd -q -w 10
```



```
$ docker run -t entrypoint_cmd -q -w 10
PING 127.0.0.1 (127.0.0.1): 56 data bytes

--- 127.0.0.1 ping statistics ---
11 packets transmitted, 10 packets received, 9% packet loss
round-trip min/avg/max = 0.069/0.138/0.634 ms
```

Any Questions?



Annex: The Layered Filesystem of Docker Images

- Overview
- Example of layering
- Image layers and container layers
- Advantages of the layered approach

Overview


- Recall the Dockerfile from the previous section, which specifies a simple image for a Python app

```
FROM python:3.8.5  
  
COPY ./main.py app.py  
  
ENTRYPOINT ["python", "app.py"]
```

- A Docker image is a layered filesystem
 - Each layer contains only the changes to the filesystem, relative to the underlying layers
- Docker uses a "union filesystem"
 - It creates a virtual filesystem for an image, out of a set of layers

Example of Layering

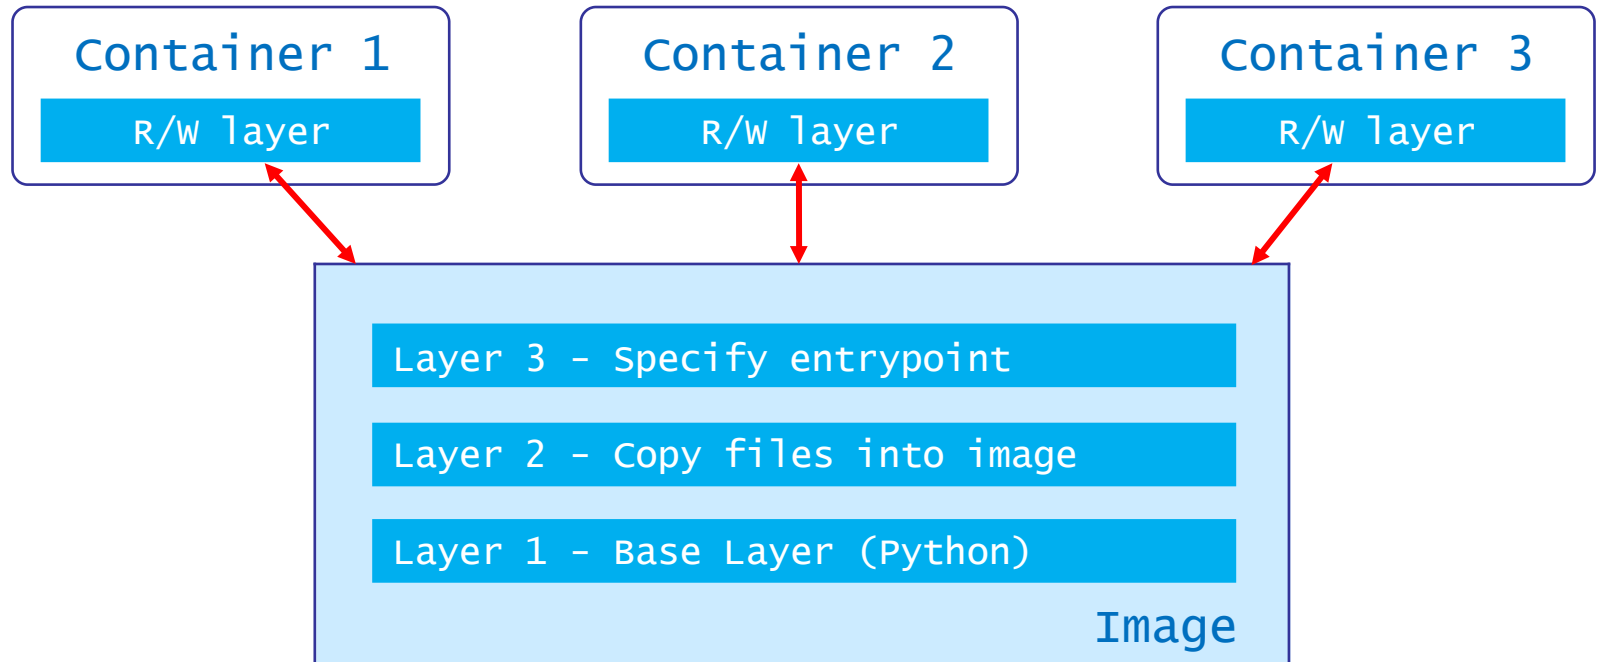
- Here are the layers for our simple Python image



- Each statement in the Dockerfile creates another layer
 - Layer 1 is the base layer (the Python distribution here)
 - Layer 2 copies files into the image
 - Layer 3 specifies the entrypoint
- The content of each layer is mapped to a special folder on the host system

Image Layers and Container Layers

- Each layer in an image is read-only
 - Layers are shared among containers created from the same image
- When Docker runs a container from an image...
 - It adds a thin R/W container layer on top, just for that container



Advantages of the Layered Approach

- Reduction of resource consumption
 - The read-only image layers are only loaded into memory once, when the first container is run
- Fast load-time for containers
 - When we run a container, Docker only needs to create the thin R/W layer for that container
 - All the underlying layers are already loaded (except for the first time we run a container)