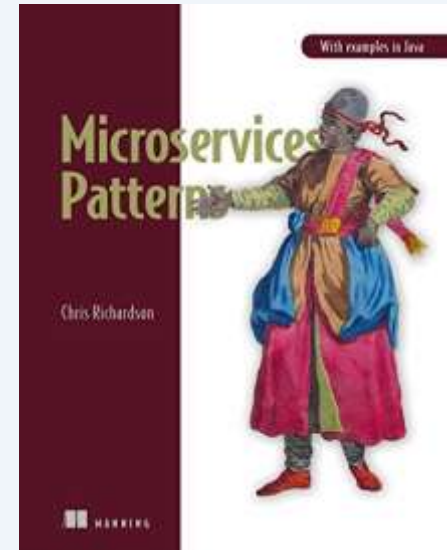


API Queries

The API Composition Pattern and CQRS

Objectives

- The need for an API query strategy
- Context & Problem
- Solution 1: The API Composition Pattern
- Solution 2: The CQRS Pattern



API Queries

- This is another problem arising from the microservices architecture pattern!
- Most applications will enable a set of queries through their API
 - getFarmerProfile()
 - findOrder()
 - findOrderHistory()
- In a monolithic architecture, these are easy to accommodate with appropriately written SQL SELECT statements to the shared relational DB.

Types of queries

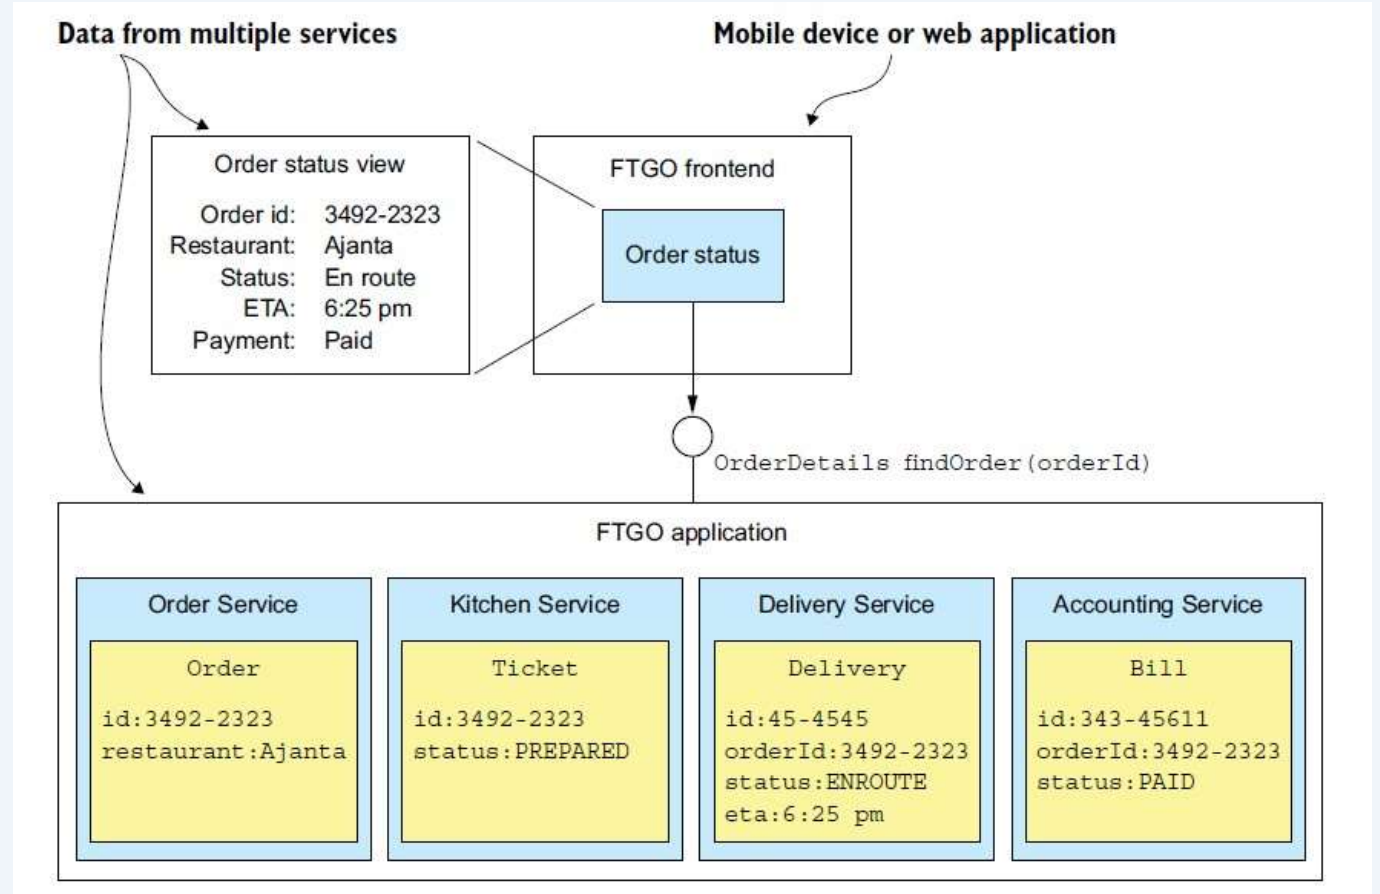
- Single service queries
 - These are easy – just a SELECT statement from the service DB
 - E.g. getFarmerProfile()
- Cross service queries
 - These are harder – the information is not stored in any one database
 - Even if it could be implemented, a distributed query would violate encapsulation.
 - E.g. findOrder(), findOrderHistory()

Context & Problem

- Your service implements the Microservices Architecture Pattern
- You have also enforced the Database per Service Pattern
- The problem is then – How to implement queries?
- The only real force involved here is the need to have a solution for implementing queries.

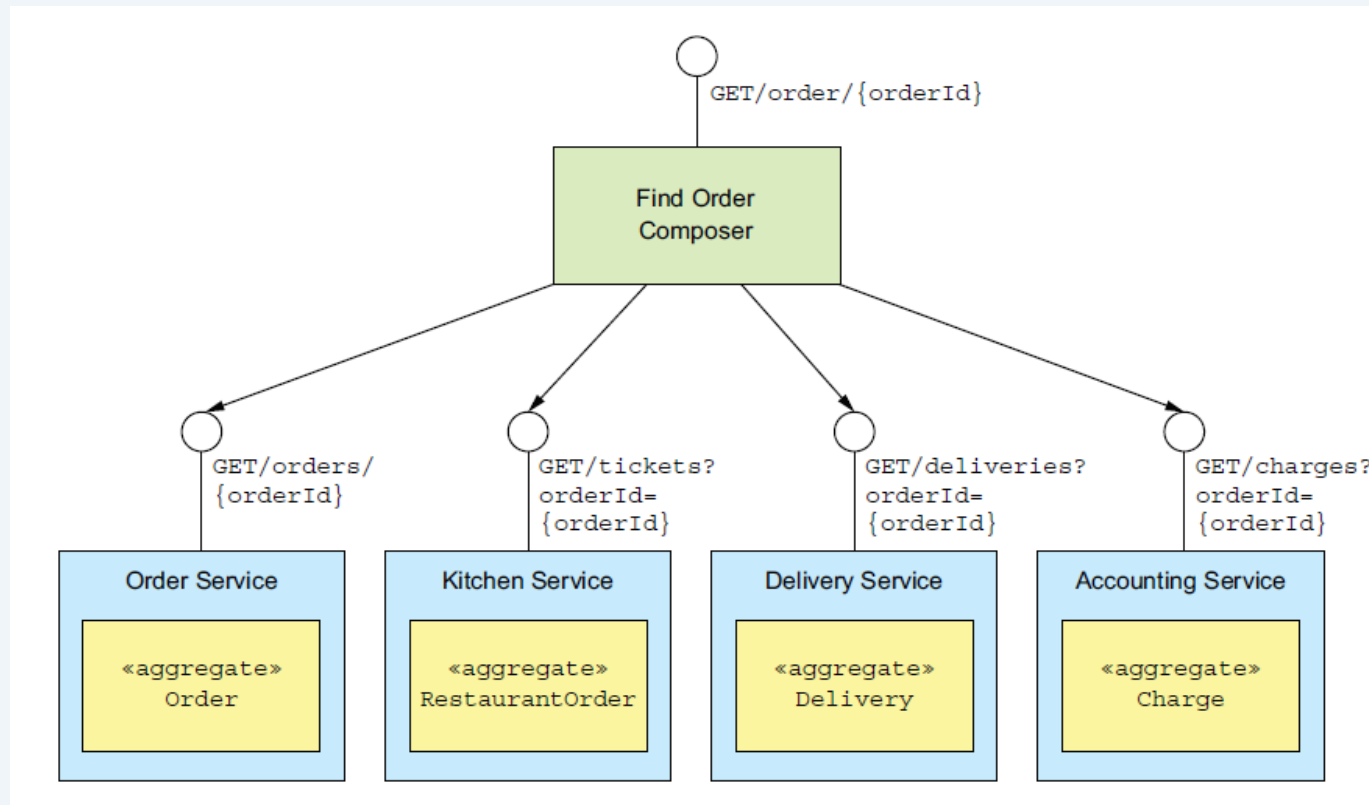
Context and Problem

- The client requires data from multiple services to complete the query



Solution 1: The API Composition Pattern

- A query is implemented by defining an **API Composer**
 - This invokes the services owning the necessary data and performs an in-memory join of the result



Example: Include with API Gateway

- An application most often has a single **API Gateway** service to route requests to the application services.
 - This often includes API composition.

Resulting Context

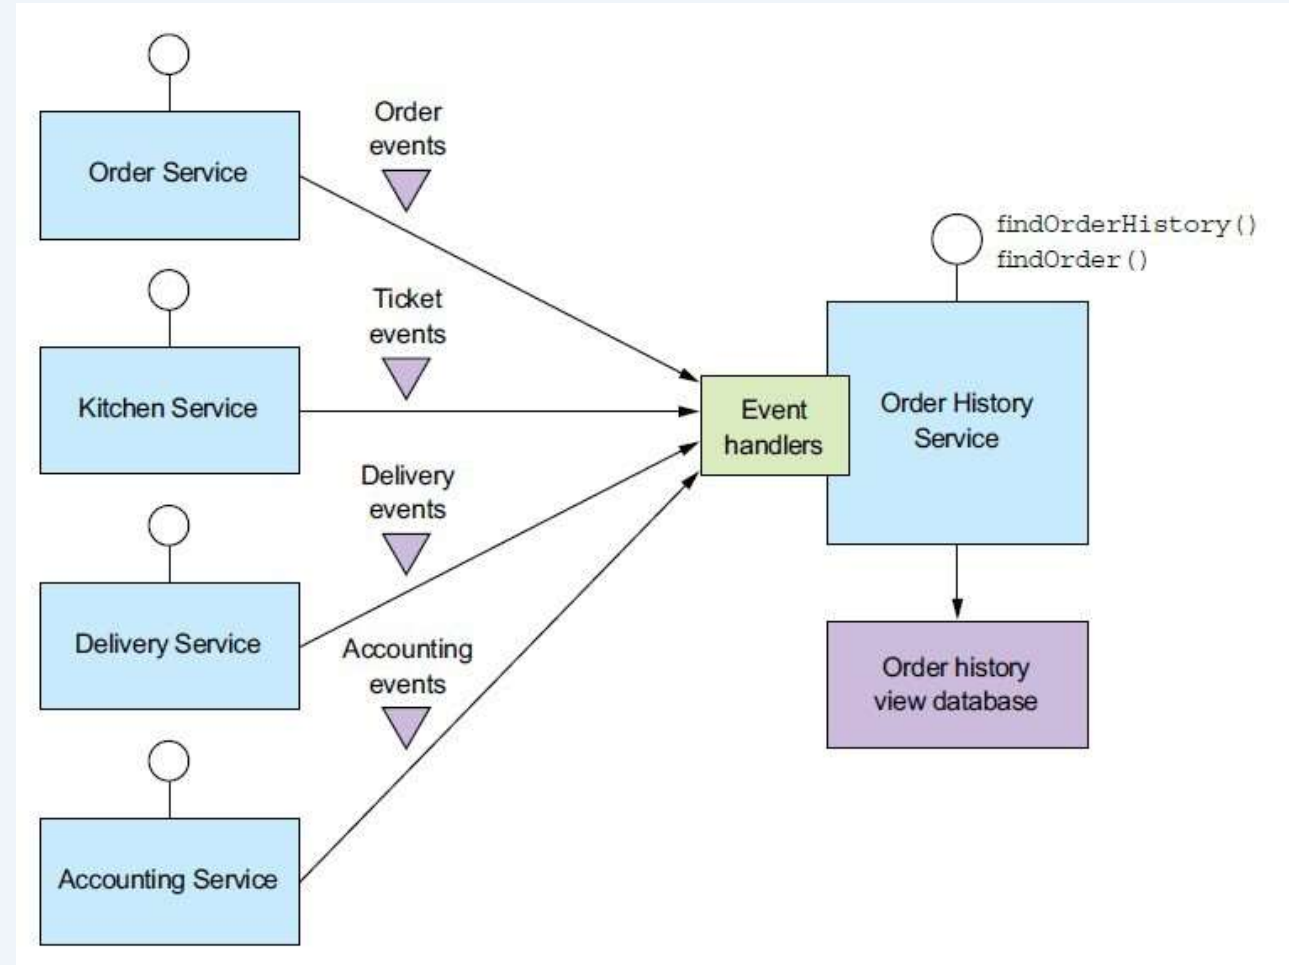
- Benefits
 - This is a very simple way to query data in a microservices architecture
- Drawbacks
 - Increased overhead – multiple services and databases queried.
 - In some cases in-memory joins of large datasets are required, which would be inefficient to implement this way → risk of reduced availability.
 - Lack of transactional data consistency
 - Because of multiple queries, there's a risk the data returned may not be consistent.
- The rule of thumb – use an API composer unless you really can't.

Solution 2: The CQRS Pattern

- Command-Query-Responsibility-Segregation
- Commands (create, update, delete) and queries (retrieve) are treated differently because the way they affect the application is different.
- The application keeps a view database, which is a read-only replica designed to support the query.
 - The replica is kept up to date by subscribing to domain events published by the service(s) that own the data.

CQRS Implementation

- This is all about separation of concerns.
- The view database doesn't belong to any particular service, so it makes sense to implement it as a standalone service.
- Replicates data that exists in other service(s), but is optimized for the particular queries allowed.



Resulting Context


- Benefits

- Enables efficient implementation of queries
- Enables efficient implementation of **diverse** queries
- Makes querying possible in an event sourcing based application
- Improves the separation of concerns

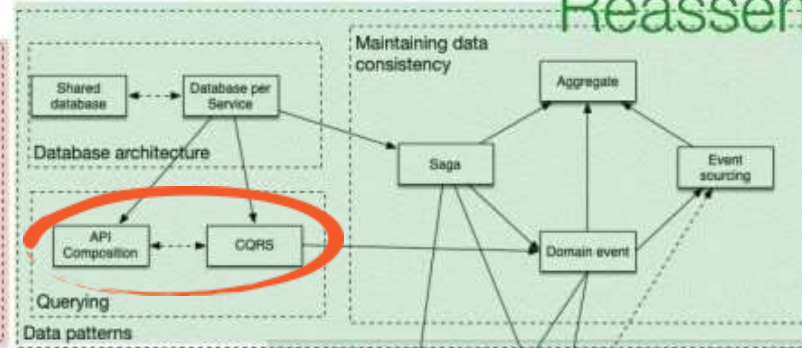

- Drawbacks

- More complex architecture
- Must deal with replication lag
 - There is a natural lag between the command-side and query-side views.
 - How does the client know if the data is out of date?
 - One solution – supply version info to the client, they can poll for updated information.



Splitting




Application patterns




Reassembling




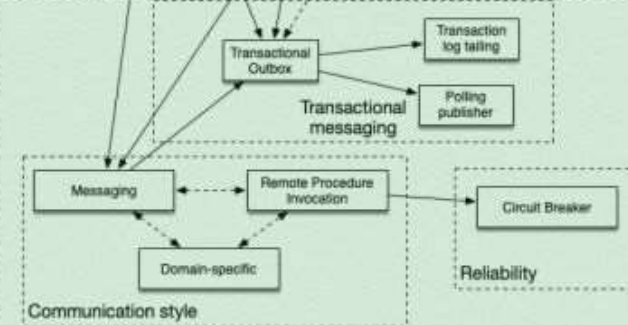
UI



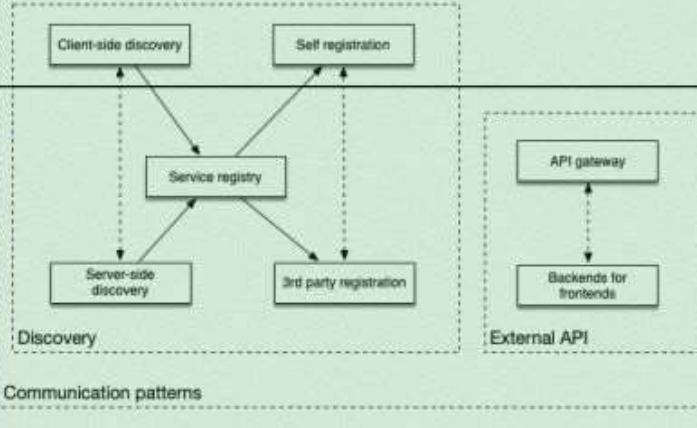

Application Infrastructure patterns



Architecture



Infrastructure patterns



Operations

Microservice patterns

Summary

- The need for an API query strategy
- Context & Problem
- Solution 1: The API Composition Pattern
- Solution 2: The CQRS Pattern

Questions or Comments?

